

FML: a VM implemented in SML

Henrik Sommerland, Oskar Ahlberg, Aleksander Lunqvist

March 5, 2014

Abstract

For our project we have decided to build a virtual machine(VM) in SML. The name FML is just an arbitrary thre letter name and has no meaning or interperatation. The VM is a RISC machine using a Von-Neuman architecture. It has a very minimalistic instruction set. The design of FML resembles those of older 8-bit architectures such as the MOS 6510 and the Z80 microprocessors commonly in use during the late 70s and early 80s. The FML machine has no “bus width” and works exclusivley with signed integersⁱ. The lack of a physical bus enables the VM to do things wich an ordinary CPU could not achive such as reading from two registers at the same time. Even though the cpu have very few operations (only 27) a very effective instruction set architecture makes these operations very flexibels and there are roughly 600 valid instruction codes. It is also noteworthy that FML is asycrounous and has now predefined clock frequenzy.ⁱⁱ

So even though FML is a very minimalistic machine it is quite powerfull. We have allso built a fully featured assembler for the FML machine.

Contents

1	Our work	2
1.0.1	Personal notes	2
2	The VM	3
2.0.2	General	3
2.0.3	Instruction Set Architecture	5
3	Assembler	6
3.0.4	General	6
3.0.5	Implementation	6

ⁱThe details of the integers used are dependent on which SML implementation is used

ⁱⁱAlthough for debugging purposues one can use both manual stepping and a fixed update speed.

4	Appendix	11
4.1	Componets.sml description	11
4.1.1	Introduction	11
4.1.2	The Ram structure	11
4.1.3	The Stack structure	12
4.1.4	The Register structure	13
4.1.5	The Program Counter structure	13

1 Our work

We have tried to work as independent as possible. This has ofcourse led to some difference in how we have commented the code, some slight differnces in naming and indentation. The way we have chosen do describe how our programs work differes a bit here depending on who wrote the code for the given part of the VM.

We have not been using any form of unit testing. But instead we have done a series of more and more complicated online tests. This due to the scale and the complexity of the various funtions and algorithms of the project.

We have continously have meetings both with just us in the group and also some with our assigned TAⁱⁱⁱ. These meetings have primairly been about informing eathother about how the VM works and how the various components of it should be implemented. We have then had an ongoing discussion on facebook regarding details and problems wich we have encountered.

We have been using Git as a source code management system. We have been using BitBucket for to host the project and troughout all of the development process the repository has been hidden and only we in the group and our TA have had access to the code.

We are using an array in the current implementation of the memory for the VM. Now i know that it is stated in the project description that the project should be written in a functional and pure way. I discussed with both Dave Clarke and Tjark Weber about using a “monad like” structure to hide the sideffects of the array hadling and they said that it was okay. The Signature handling the memory is written in such a way that any other part of the program implementing the memmory structure will not be able to se that there are any side effects. I.e there are no semantically observable side effects of the memory structure. All of the code would look exactly the same from “the outside” regardless of how we implemented the memory.

1.0.1 Personal notes

In this section we will give som personal notes regarding our parts in the project.

ⁱⁱⁱOur TA has during this project been Tobias Neil

1.0.1.1 Henrik Sommerland I have been incharge of designing the VM and writing the assembler. I have also written some signatures for the others in the group to help them get started. I began work very early, as soon as we had gotten permission to start on the project. I began by writing the specifications for the VM.

Even though i have never had any formal education in how computer architectures work I have learnt a lot about it on my own. When i was younger (around 17) i designed a 8-bit cpu from TTL logic chips (the 7400 series). It was during this time I learned how to build a cpu. So the design of the FML machine was for me very straight forward. I have done all of the design myself and i have not copied anything from books or any previous designs. Although my way of thinking and reasoning about cpu architectures comes primarily from my own work and the designs of older 8-bit cpus. The design of the cpu took roughly one or two days to finish and then there has been a continous process of ironing out bugs and inconsitencies.

Then I started to write the assembler. From the start I had a pretty good idea about what I wanted the assembler to do and I had a pretty good idea how to implement it. I wrote the assembler in about three days and I then had a fully working assembler.

After I had completed the assembler I sat out to start testing it and to write the documentation for it. As I wrote the documentation and did more extensive testing of the assembler I worked out the last few bugs and such in the code.

Then i started to write signature files and such for the others in the group to get to work on. I allso started to write on the major report for the entire program whilst guiding the others in the group.

In general I have found this project to be incredibly fun and interesting. In the beggining I was worried that the others in the group where not going to be able to understand how it all woked and even though it's really hard to explain something complex wich is crystal clear in my head to other people the others in the group have ben very enthusiastic and have really pulled trough this mammoth of a project.

2 The VM

Here a informal description of the workings of the machine. For a more detailed description se the VM specifications in the appendix.

2.0.2 General

The FML machine is built up as a very simple von-neuman architecture. The machine consists ony of a few major components. It's notewhorthy that there is no instruction decoder present. This is since all of the instruction decoding and handling takes place within the software implementation of the machine. The size of the memory the machine has avaiable is arbitrary and is defined

at the initialization of the machine. Below will follow a dataflow diagram of the machine, describing all of the components and how they can communicate.

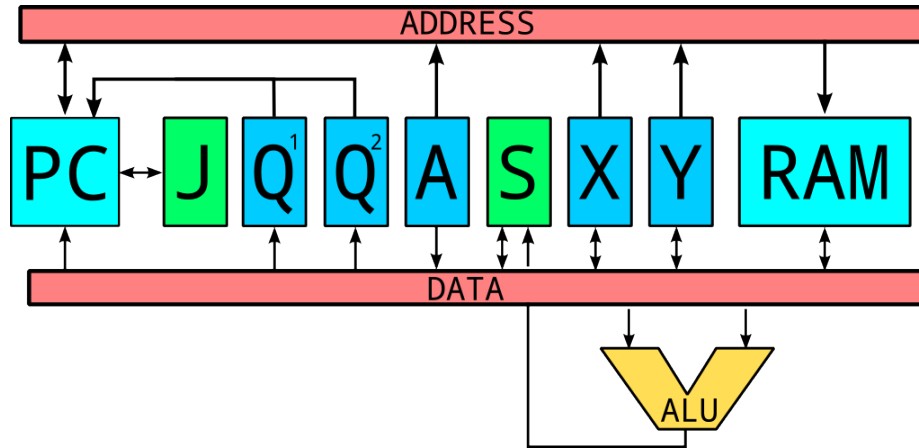


Figure 1: Dataflow diagram of the FML machine

Now this image might be a little bit confusing. One should consider the two read rectangles DATA and ADDRESS as “virtual buses”. One can interpret the picture as: X can both read and write from other components and be used for addressing. Below will follow brief descriptions of the components. More indepth descriptions are given in the appendix.

X and Y

These are the two general purpose registers wich can be read, written to and used for addressing.

S

This is the general purpose stack. It can be both read and written to. Ev-erytime some thing gets written to the stack it gets pushed onto the stack and everytime something is read from the stack the stack gets popped. The stack can not be used for addressing.

A

This is only a virtual register. It is read only and can be used for address-ing. This is only used if an instructuion uses a non-register argument^{iv}.

Q¹ and Q²

These are the two interupt registers. These are very special and can only be written to. They will hold the addresses to wich the machine should jump if an prepherial component makes a interupt request. More on this later.

^{iv}A non-registry argument is a argument wich is not any of the registers, the stack, or something from the memmory. The value of A will (if used) be at the memmory cell directly folowing the one at wich the program counter is.

PC

This is the program counter. It keeps track on where in the memory the instructions are being read from. It also handles the jumping.

J

This is the jump stack. This stack is used to store the return addresses for subroutine jumps. This stack can only be manipulated by the program counter.

ALU

This is not really a ALU. The machine does not have a separate ALU component but this is just here to illustrate that the all of the components which can be read from can be used as arguments for arithmetic and logical operations. All of the results from the arithmetic and logical operations are always put on the stack.

RAM

This is the random access memory of the machine.

2.0.3 Instruction Set Architecture

The ISA of the VM is built in a special but simple fashion. Each instruction corresponds to a six digit integer.. Where each digit corresponds to specific information regarding different types of opcodes. The digits counting from right to left is:

First Second argument

Second First argument

Third Arithmetic operations

Fourth Logic operations

Fifth Jump operations

Sixth Special

This system of encoding information into each digit of the instruction makes the implementation of the instruction controller and the assembler much easier. It allows for all the operation types to be grouped into numerical ranges and it gives a lot of flexibility. Note that some of the instructions may be invalid and some might be nonsensical but the instruction controller crashes if a invalid instruction is encountered. The assembler is written in such a way that it can only generate valid instructions. So an example would be: 000401. Where the 4 tells us that it we should perform a modulo operation, the 0 says that the second argument is the **X**register and the last 1 says that the first argument is the **Y**register. Notice that the order of the last two digits is reversed in respect to the order of the arguments in the operation. This is due to a design choice made

early in the design phase. It makes the instructions code a bit more confusing to read but it makes the assembly code become far more intuitive. For a more detailed description of the ISA see the VM specifications in the appendix.

3 Assembler

3.0.4 General

The assembler which we have written for the FMI machine is a very basic yet powerful assembler. The assembler doesn't do much more than catch invalid opcodes and arguments. It also enables the use of both label pointers and value pointers. The main tasks of the assembler are the instruction code generation and address resolution. The syntax of the assembler is inspired by the syntax for the MOS 6510 assembly language and primarily the syntax of the Turbo Assembler for the Commodore 64. The assembler is now fully functional and there we don't see any need to augment it or redesigning any aspects of it. The assembler should only generate valid instructions but due to a major design error in the implementation of the VM this is not necessarily true any more. Below a short example of an assembly program will follow:

```
% This is a simple program which fills a part
% of the memory with 100 consecutive integers
% through relative addressing.
#start
MOV 0 x
@start_address
MOV start_address y
#loop
MOV x $y
INC x
INC y
BLE x 100
JMP loop
HLT
```

A line starting with a `#` declares a label. The address of the label will correspond to where in the code the label gets declared. A line starting with a `@` declares a value. The address of the label will be assigned independent of where in the code it appears. For a more in-depth description of how the assembly language see the Assembler part of the description.

3.0.5 Implementation

3.0.5.1 General The assembler works in a fairly straightforward way. The first step in the process of assembling is the lexical analysis in which

the lines in the text file gets tokenized. In this stage an “intermediate structure”^v gets constructed. This is an object which contains all of the labels, values and a list of the tokenized lines. The list of the tokens contains tuples of (`label`,`offset`,`token`) where the label is the last initialized label and the offset is how many addresses away from that line the current token is. All of the labels and values will not be assigned an address in this phase. It is in this phase where the opcodes and their arguments get converted into their corresponding numerical instruction code. It is also during this phase in which the syntax gets checked. If a syntax error is encountered the assembler will stop immediately. When the lexicographical analysis has been completed a check for duplicate pointer declaration is performed.

The next phase in the assembly is the address resolution phase. This is done in two phases, in the first one the labels get resolved and in the second the values get resolved. It begins by first resolving the labels. This is done by first giving the assembler a *base address* which is the address of the first label. As of now the first non comment line in the input file has to be a label since every line has to have a label assigned to it. Then the address resolution function continues down the intermediate structure and remembering which line it is at and what its last read label was. When it runs into a new label token it will set the new label to its current address and then continue on until it has gone through the entire intermediate structure. After the labels have been resolved the assembler starts to resolve the values. This is done in a very straightforward way. The assembler just looks at the last address of the last entry in the output of the first pass and looks at the last address, adds one to it and then just places all the values in after that address in the same order as they appeared in the file. After all the address has been resolved the assembler runs through the list of tokens and replaces every pointer token with its correct address.

After this is completed the assembler finalizes the code by converting everything into a list of integers which then gets outputted to a file. And that children is how assembly code gets turned in to machine code.

the assembler runs in linear time with respect to the number of lines in the code. This is under the assumption that the number of lines are far greater than the number of values and labels in the code. This is a safe assumption for any reasonably written code.

3.0.5.2 Flow chart Below a flow chart will follow for how the assembler assembles the assembly code.

^vThe use of the word structure here is a bit ambiguous since it actually is a structure in sml. But in this text it will refer to an abstract structure of data.

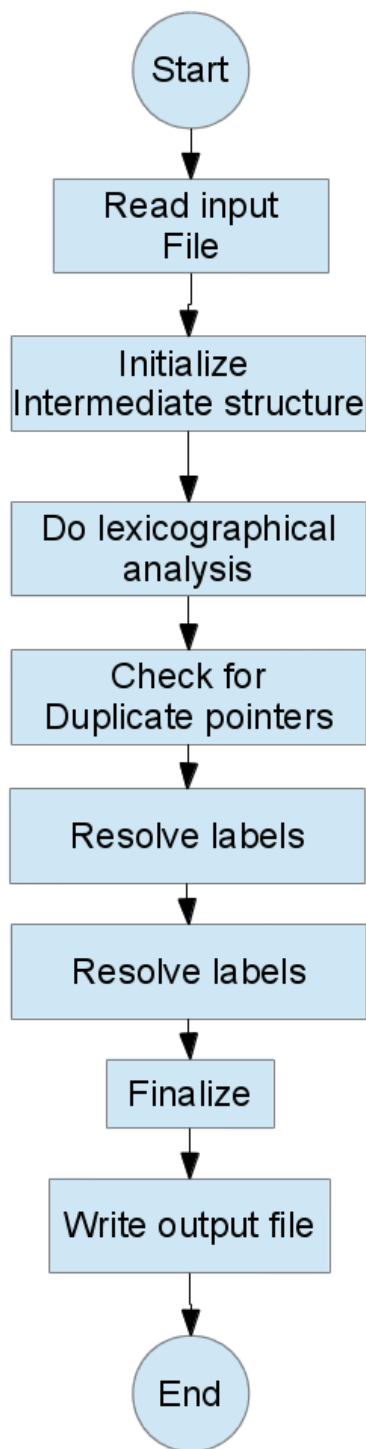


Figure 2: Dataflow diagram of the assembler

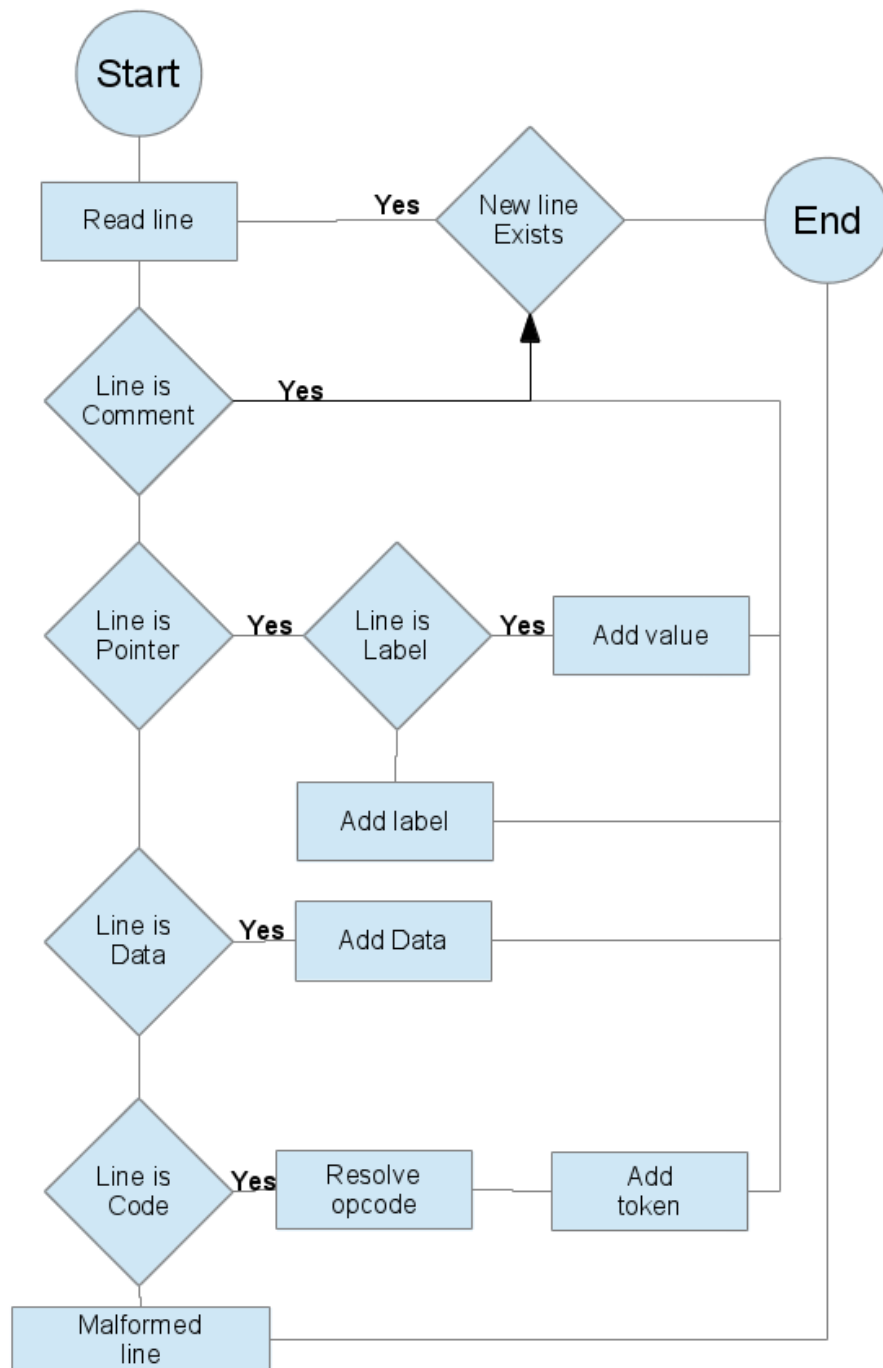


Figure 3: Dataflow diagram of the tokenization phase




Figure 4: Dataflow diagram of the label address resolution

The i have not included a flowchart for the address resolution of the values or the finalization part since these are trivial.

3.0.5.3 Usage To use the assembler properly one has to know how to write assembly code and understand the detailed workings of the machine. We recomend studying both the VM spcifications and the assembler documentation in the appendix before you start to write programs for the machine.

The working of the assembler program is very straight forward. just frite your assembly code in a file called `in.asm` and run the `Assembler.sml` file in the sml interperter of your choosing and if there are no errors encountered during the assembling of the program the assembeled program will be outputted to a file named `out.fml`.

4 Appendix

4.1 Componets.sml description

Due to misscommunication a radically diffrent description of how the Components.sml file works was written and have been included here as an appendix.

4.1.1 Introduction

The following structures and signatures are present in Components.sml are the Ram, Stack, Register and ProgramCounter.

4.1.2 The Ram structure

4.1.2.1 Synopsis

signature RAM
structure Ram :>RAM

The Ram structure provides a base of the functions of a ram memory. This structure acts as something akin to a “monad”. It hides all of the sidefects sued for the array handling.

4.1.2.2 INTERFACE

```
type memory = int array
val initialize : int → memory
val getSize : (memory) → int
val write :(memory * int * int ) → memory
val read : (memory * int) → int
val load : (memory* int list) → memory
val writeChunk : (memory * int * (int array)) → memory
val readChunk : (memory * int * int) → int array
val dump : memory → string
```

4.1.2.3 Description

val initialize : int \rightarrow memory

Initialize the ram to a memory with the size of int, when int \neq 0

val getSize : (memory) \rightarrow int

Gets the size of the memory

val write : (memory * int * int) \rightarrow memory

write takes a memory and writes a new value of int at the pointer of the first int and returns the memory

val read : (memory * int) \rightarrow memory

read takes a memory and reads the value of the place of int

val load : (memory * int list) \rightarrow memory

load takes a list of values and loads them to the memory

val writeChunk : (memory * int * (int array)) \rightarrow memory

writeChunk takes a memory and a start pointer and adds a chunk to the memory

val readChunk : (memory * int * int) \rightarrow int array

readChunk takes a memory and reads a chunk from first int to the last int and gives the values as an int array

val dump : memory \rightarrow string

dump takes a memory and returns the value as strings

4.1.3 The Stack structure

4.1.3.1 Synopsis

signature STACK

structure Stack :>STACK

The Stack structure provides a base for the stack part of the Pc structure.

4.1.3.2 INTERFACE

datatype stack = Stack of (int list)

val empty : stack

val push : stack * int \rightarrow stack

val pop : stack \rightarrow stack

val top : stack \rightarrow int

val isEmpty : stack \rightarrow bool

val dumpStack : stack \rightarrow string

4.1.3.3 Description

val empty : stack

is a definition of a empty Stack

val push : stack * int \rightarrow stack

takes a stack and adds the value of int to the stack.

val pop : stack \rightarrow stack

takes a Stack and pops the first element of the stack.

val top : stack \rightarrow int

takes the stack and returns the first element of the stack

```

val isEmpty : stack → bool
takes a stack and checks if it is empty if it is then true else false.
val dumpStack : stack → string
takes a stack, then pops the stack until its empty and returns all values as string

```

4.1.4 The Register structure

4.1.4.1 Synopsis

```

signature REGISTER
structure Register :>REGISTER

```

The Register structure provides a base structure of the different register that is contained in the Pc as well the Virtual machine. The vm has two different registers.

4.1.4.2 INTERFACE

```

datatype reg = Reg of int
val setData : (reg * int) → reg
val getData : reg → int
val increment : reg → reg
val decrement : reg → reg
val dumpRegister : reg → string

```

4.1.4.3 Description

```

    val setData : (reg * int) → reg
Setups a new Register
val getData : reg → int
Gets the value of the reg as an int
val increment : reg → reg
Takes a reg and increment it with one.
val decrement : reg → reg
Takes a reg and decrements it with one.
val dumpRegister : reg → string
Takes the register and adds all elements to a string.

```

4.1.5 The Program Counter structure

4.1.5.1 Synopsis

```

signature PROGRAM_COUNTER
structure ProgramCounter :>PROGRAM_COUNTER

```

The ProgramCounter structure controls the execution flow of the VM

4.1.5.2 INTERFACE

```
datatype pc = Pc of (int * Stack.stack * Register.reg * Register.reg)
val incrementPointer : (pc * int) → pc
val jump : (pc * int) → pc
val subroutineJump : (pc * int) → pc
val return : pc → pc
val interrupt : (pc * int) → pc
val dumpPc : pc → string
```

4.1.5.3 Description

```
val incrementPointer : (pc * int) → pc
Takes a Pc and adds a int  $i$  0
val jump : (pc * int) → pc
Takes a Pc and jumps the pc counter to the value of int  $i$  0
val subroutineJump : (pc * int) → pc
Takes a Pc and preforms SubrutineJump with the value of int  $i$  0 and adds the
value of the pointer + 1 to the stack
val return : pc → pc
Takes a pc and gets the value from the pointer and pops the stack with the
value
val interrupt : (pc * int) → pc
if the value of a is 1 or 2, then the value of i is added to s
val dumpPc : pc → string
Takes a pc and dumps the content of the pc as a string (the Pc contained a
pointer, Stack, and tow registers)
```